

When I started the General AI fluency internship I thought I'd just be building a portfolio and learning how to prompt better. Disclaimer: It was much more than that. I built my portfolio, yeah, and learnt how to prompt better, but I also made my own AI agent, publicized the portfolio, completed the UX/UI design internship alongside this one, and am now writing this as part of the last assignment for this track.

Nothing changed except my knowledge, focus, and ability to hold myself accountable to deadlines. Which basically means everything changed.

One of the biggest changes was getting a closer look at how AI systems are actually structured. What looks like a tiny action to a user can involve a surprising number of nodes, conditions, API calls, guardrails, limitations and human judgement on the backend. Building my tutoring agent made that especially clear. Something that sounds as simple as *"check whether a student has joined and contact the parent if they haven't"* became a workflow involving Google Calendar, Gmail, Zoom, stored parent information, attendance checks, timers, parent-response detection and different levels of human approval. I'm walking away with a newfound respect for AI engineers and developers.

There are three things in particular that I'm taking away from the experience.

1. Start small. Focus on the functional requirements before considering the non-functional ones, and ensure quality over quantity. It is very easy to imagine everything a system *could* do before making sure its core functionality actually works. Building incrementally made it easier to test individual parts, find problems and understand what was actually happening instead of creating a huge workflow that was difficult to debug.

2. AI isn't bad — use it for collaboration, not replacement. Before this internship, I deliberately avoided using AI for a lot of my work because I was afraid it would strangle my thinking skills and learning. I've seen that happen to more people around me than I care to admit. I still think that concern is valid, but my perspective on how I personally use AI has changed. I can use it to question my decisions, troubleshoot problems, explore alternatives and speed up repetitive parts of my work without handing over the thinking itself. That perspective feels much more productive, matches where the market is going, and still allows me to stay in charge of the work.

3. Consider the blast radius. This was a new term for me, but a very accurate one. Not every AI action carries the same consequences when it goes wrong. That became important when deciding what my agent could do automatically and what should require my approval. A simple parent check-in has a relatively small blast radius. Changing a student's schedule or sending a final no-show communication that mentions the possibility of charging for a session is different. Those decisions stayed behind human approval. I've learnt to consider the consequences of an incorrect action, not just whether that action *can* be automated.

If I continued developing the agent, my next focus would actually be less about adding features and more about making the existing system more efficient. I would reduce

unnecessary computational and API costs. For example, the monitoring agent currently retrieves Calendar events and then filters them. I would investigate filtering the events before fetching them so that irrelevant events are never retrieved in the first place. I would also redesign the monitoring trigger so that it does not continue checking the Calendar for new sessions every five minutes throughout an existing session regardless of whether the student is already attending. That would reduce unnecessary requests and processing.

If I could say one thing to the version of myself who started in Week 1, it would probably be that learning AI fluency isn't really about learning how to get AI to do more of your work. For me, it became much more about learning where AI is useful, where I need to stay involved, and how to tell the difference.